

# Details about Midterm Exam

Date & Time: **Thursday, October 7**

Check your email for the time you are scheduled to take the midterm exam. Let me know if there are conflicts well in advance.

- **All exams will be proctored with Honorlock.**

In-class review: **Monday, October 4**

Post questions in the [Midterm Exam Discussion Board](#) until **Sunday, October 3 @ 3 PM**.

- These questions will be addressed in-class on Monday, October 4.

## Coverage

Lectures 1-17 (modules 1-4)

## Practice Midterm Exam

1. Available in the [Assignment-Solutions](#) repo.
2. Practice Canvas Quizzes created to practice Honorlock proctoring and overall exam setup: [Practice Midterm Exam - Honorlock Proctoring](#)
3. Solutions to the practice exam will be posted before the in-class midterm review.
4. One of the questions in the practice exam asks about K-Means clustering, but K-Means clustering will **not** be covered in the Midterm exam.

## Allowed Material for the Midterm

- 1-page letter-sized of **formulas** (front and back, handwritten or typed)
- scientific calculator
- Writing and scratch paper to write your answers. I will give an estimate on how many pages to bring in the review lecture.
- Bring a second device (e.g., phone, tablet or other) with the [CamScanner](#) or [Scannable](#) app installed. This device is only to be used at the end of the exam to take pictures of your handwritten solutions
- **TOTAL TIME:** 2 hours + 15 minutes

**Communications between students or anyone else during the exam is considered cheating.  
Turn off all Slack notifications and other communications channels!**

Ten (10) minutes prior to the time you are scheduled to take the exam, you will be able to see the Canvas quiz with your midterm exam.

---

# Lecture 15 - Naive Bayes Classifier; Mixture Models

## Naive Bayes Classifier

Therefore, for a given test point  $\mathbf{x}^*$ , our decision rule is:

$$P(C_1|\mathbf{x}^*) \underset{C_2}{\overset{C_1}{\geq}} P(C_2|\mathbf{x}^*)$$

Using the Bayes' rule, we can further rewrite it as:

$$\frac{P(\mathbf{x}^*|C_1)P(C_1)}{P(\mathbf{x}^*)} \underset{C_2}{\overset{C_1}{\geq}} \frac{P(\mathbf{x}^*|C_2)P(C_2)}{P(\mathbf{x}^*)}$$

This defines the **Naive Bayes Classifier**.

## Training a Generative Classifier

So, **to train the classifier**, what we need to do is to determine the parametric form (and its parameters values) for  $p(x|C_1)$ ,  $p(x|C_2)$ ,  $P(C_1)$  and  $p(C_2)$ .

- For example, we can assume that the data samples coming from either  $C_1$  and  $C_2$  are distributed according to Gaussian distributions. In this case,

$$p(x|C_k) = \frac{1}{(2\pi)^{1/2}|\Sigma_k|^{1/2}} \exp\left\{-\frac{1}{2}(\mathbf{x} - \mu_k)^T \Sigma_k^{-1}(\mathbf{x} - \mu_k)\right\}, \forall k = \{1, 2\}$$

For this case, this defines the **Gaussian Naive Bayes Classifier**.

Or, we can consider any other parametric distribution.

- What about the  $P(C_1)$  and  $P(C_2)$ ?

We can consider the relative frequency of each class, that is,  $P(C_k) = \frac{N_k}{N}$ , where  $N_k$  is the number of points in class  $C_k$  and  $N$  is the total number of samples.

## MLE Parameter Estimation Steps

For simplification, let's consider the covariance matrix  $\Sigma_k$  for  $k = 1, 2$  to be **isotropic** matrices, that is, the covariance matrix is diagonal and the element along the diagonal is the same, or:  $\Sigma_k = \sigma_k^2 \mathbf{I}$ .

- What are the parameters? The mean and covariance of the Gaussian distribution for both classes.

Given the assumption of the Gaussian form, how would you estimate the parameters for  $p(x|C_1)$  and  $p(x|C_2)$ ? We can use **maximum likelihood estimate** for the mean and covariance, because we are looking for the parameters of the distributions that *maximize* the data likelihood!

The MLE estimate for the mean of class  $C_k$  is:

$$\mu_{k,\text{MLE}} = \frac{1}{N_k} \sum_{n \in C_k} \mathbf{x}_n$$

where  $N_k$  is the number of training data points that belong to class  $C_k$ .

Assuming the classes follow a (bivariate or 2-D) Gaussian distribution and, for simplicity, let's assume the covariance matrices are **isotropic**, that is,  $\Sigma_k = \sigma_k^2 \mathbf{I}$ .

The MLE steps for parameter estimation are:

1. Write down the observed data likelihood,  $\mathcal{L}^0$

$$\mathcal{L}^0 = P(x_1, x_2, \dots, x_N | C_k) \quad (1)$$

$$= \prod_{n=1}^N P(x_n | C_k), \text{ data samples are i.i.d.} \quad (2)$$

$$= \prod_{n=1}^N \frac{1}{(2\pi)^{1/2} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (x_n - \mu_k)^T \Sigma_k^{-1} (x_n - \mu_k) \right\} \quad (3)$$

$$= \prod_{n=1}^N \frac{1}{(2\pi)^{1/2} |\sigma_k^2 \mathbf{I}|^{1/2}} \exp \left\{ -\frac{1}{2\sigma_k^2} (x_n - \mu_k)^T \mathbf{I} (x_n - \mu_k) \right\} \quad (4)$$

$$= \prod_{n=1}^N \frac{1}{(2\pi)^{1/2} (\sigma_k^2)^{1/2}} \exp \left\{ -\frac{1}{2\sigma_k^2} (x_n - \mu_k)^T (x_n - \mu_k) \right\} \quad (5)$$

1. Take the log-likelihood,  $\mathcal{L}$ . This *trick* helps in taking derivatives.

$$\mathcal{L} = \ln(\mathcal{L}^0) \quad (6)$$

$$= \sum_{n=1}^N -\frac{1}{2} \ln 2\pi - \frac{1}{2} \ln \sigma_k^2 - \frac{1}{2\sigma_k^2} (x_n - \mu_k)^T (x_n - \mu_k) \quad (7)$$

1. Take the derivative of the log-likelihood function with respect to the parameters of interest. For Gaussian distribution they are the mean and covariance.

$$\frac{\partial \mathcal{L}}{\partial \mu_k} = 0 \quad (8)$$

$$\sum_{n \in C_k} \frac{1}{\sigma_k^2} (x_n - \mu_k) = 0 \quad (9)$$

$$\sum_{n \in C_k} (x_n - \mu_k) = 0 \quad (10)$$

$$\sum_{n \in C_k} x_n - \sum_{n \in C_k} \mu_k = 0 \quad (11)$$

$$\sum_{n \in C_k} x_n - N_k \mu_k = 0 \quad (12)$$

$$\mu_k = \frac{1}{N_k} \sum_{n \in C_k} x_n \quad (13)$$

This is the sample mean for each class. And,

$$\frac{\partial \mathcal{L}}{\partial \sigma_k^2} = 0 \quad (14)$$

$$\sum_{n \in C_k} -\frac{1}{2\sigma_k^2} + \frac{2(x_n - \mu_k)^T (x_n - \mu_k)}{(2\sigma_k^2)^2} = 0 \quad (15)$$

$$\sum_{n \in C_k} -1 + \frac{(x_n - \mu_k)^T (x_n - \mu_k)}{\sigma_k^2} = 0 \quad (16)$$

$$\sum_{n \in C_k} \frac{(x_n - \mu_k)^T (x_n - \mu_k)}{\sigma_k^2} = N_k \quad (17)$$

$$\sigma_k^2 = \sum_{n \in C_k} \frac{(x_n - \mu_k)^T (x_n - \mu_k)}{N_k} \quad (18)$$

This is the sample variance for each class. Then we can create  $\Sigma_k = \sigma_k^2 \mathbf{I}$ , which is the (biased) sample covariance for each class.

In practice, if we want to estimate an entire covariance matrix, we would have to take the derivative of the log-likelihood function with respect to every entry in the covariance matrix.

We can determine the values for  $p(C_1)$  and  $p(C_2)$  from the number of data points in each class:

$$p(C_k) = \frac{N_k}{N}$$

where  $N$  is the total number of data points.

```
In [1]: import numpy as np
from scipy.stats import multivariate_normal
import matplotlib.pyplot as plt
plt.style.use('bmh')
```

```
In [2]: def generateData(mean1, mean2, cov1, cov2, N1, N2):
# We are generating data from two Gaussians to represent two classes
# In practice, we would not do this - we would just have data from the problem we a
```

```

data_C1 = np.random.multivariate_normal(mean1, cov1, N1)
data_C2 = np.random.multivariate_normal(mean2, cov2, N2)

plt.scatter(data_C1[:,0], data_C1[:,1], c='r')
plt.scatter(data_C2[:,0], data_C2[:,1])
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.axis([-4,4,-4,4])
return data_C1, data_C2

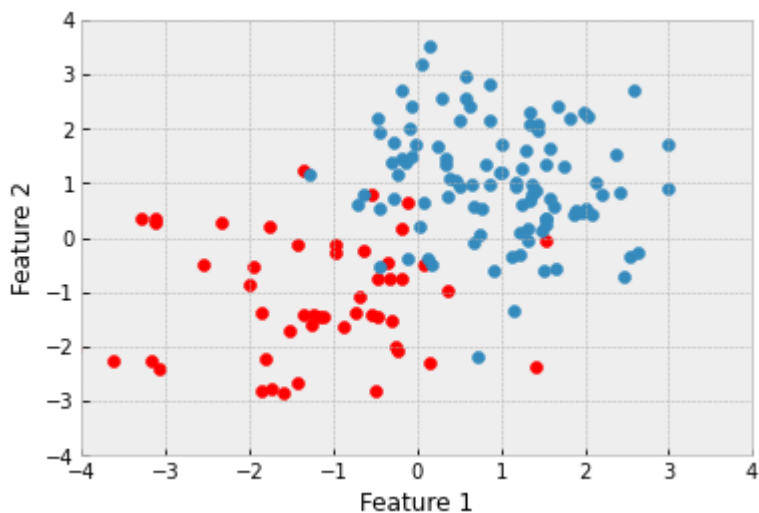
```

```

In [13]: mean1 = [-1, -1]
mean2 = [1, 1]
cov1 = [[1,0],[0,1]]
cov2 = [[1,0],[0,1]]
N1 = 50
N2 = 100

data_C1, data_C2 = generateData(mean1, mean2, cov1, cov2, N1, N2)

```



```

In [14]: data_C1.shape

```

```

Out[14]: (50, 2)

```

```

In [15]: # Estimate the mean and covariance for each class from the training data

mu1 = np.mean(data_C1, axis=0)
print('Mean of Class 1: ', mu1)

cov1 = np.cov(data_C1.T) # because np.cov assumes data comes in as DxN
print('Covariance of Class 1: ', cov1)

mu2 = np.mean(data_C2, axis=0)
print('Mean of Class 2: ', mu2)

cov2 = np.cov(data_C2.T)
print('Covariance of Class 2: ', cov2)

```

```
# Estimate the prior for each class

pC1 = data_C1.shape[0]/(data_C1.shape[0]+data_C2.shape[0])
print('Probability of Class 1: ',pC1)

pC2 = data_C2.shape[0]/(data_C1.shape[0]+data_C2.shape[0])
print('Probability of Class 2: ',pC2)
```

```
Mean of Class 1: [-1.21519443 -1.10616409]
Covariance of Class 1: [[1.44725922 0.06498299]
 [0.06498299 1.16522598]]
Mean of Class 2: [0.95239041 1.03401998]
Covariance of Class 2: [[ 0.83583401 -0.13410859]
 [-0.13410859 1.08745696]]
Probability of Class 1: 0.3333333333333333
Probability of Class 2: 0.6666666666666666
```

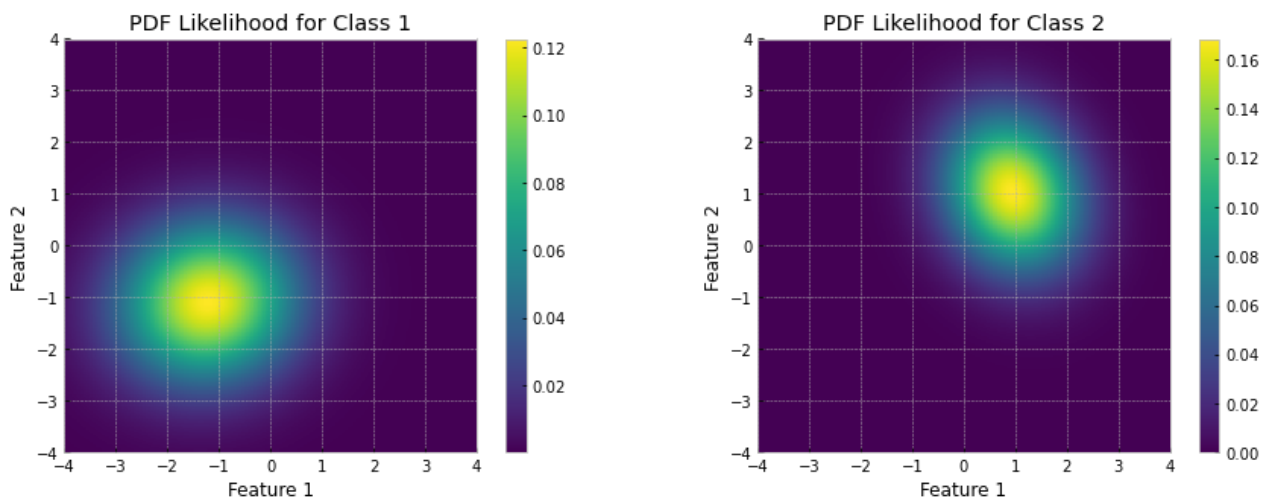
In [16]:

```
# Compute a grid of values for x and y
x = np.linspace(-4, 4, 100)
y = np.linspace(-4, 4, 100)
xm, ym = np.meshgrid(x, y)
X = np.flip(np.dstack([xm,ym]),axis=0)

# Let's plot the probabaility density function (pdf) for each class
y1 = multivariate_normal.pdf(X, mean=mu1, cov=cov1) # P(x|C1)
y2 = multivariate_normal.pdf(X, mean=mu2, cov=cov2) # P(x|C2)

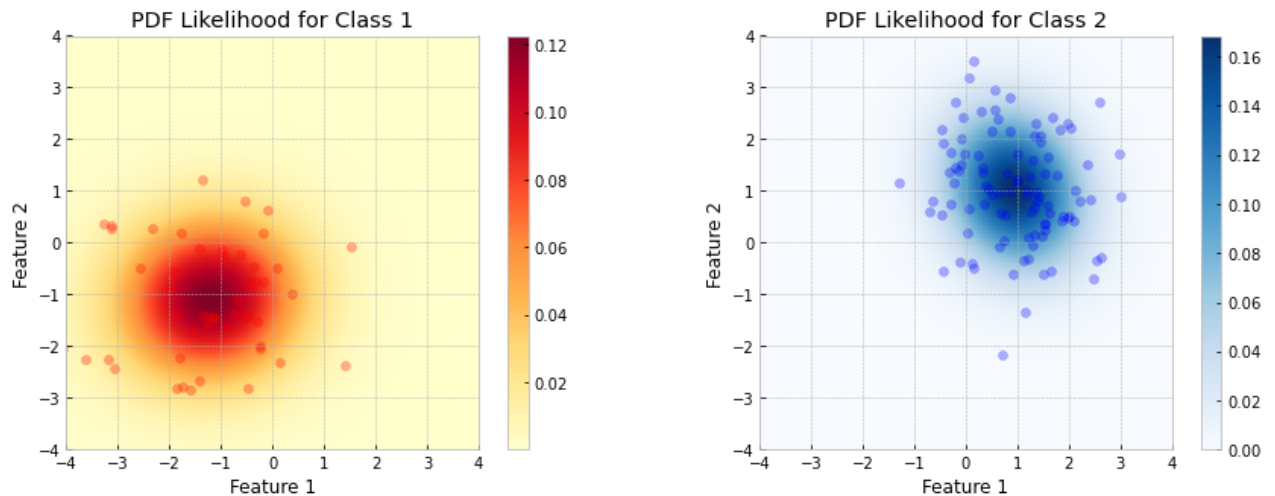
fig =plt.figure(figsize=(15,5))
fig.add_subplot(1,2,1)
plt.imshow(y1, extent=[-4,4,-4,4])
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('PDF Likelihood for Class 1')

fig.add_subplot(1,2,2)
plt.imshow(y2, extent=[-4,4,-4,4])
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('PDF Likelihood for Class 2');
```



```
In [17]: fig=plt.figure(figsize=(15,5))
fig.add_subplot(1,2,1)
plt.scatter(data_C1[:,0], data_C1[:,1], c='r',alpha=0.3)
plt.imshow(y1, extent=[-4,4,-4,4],cmap='YlOrRd')
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('PDF Likelihood for Class 1')

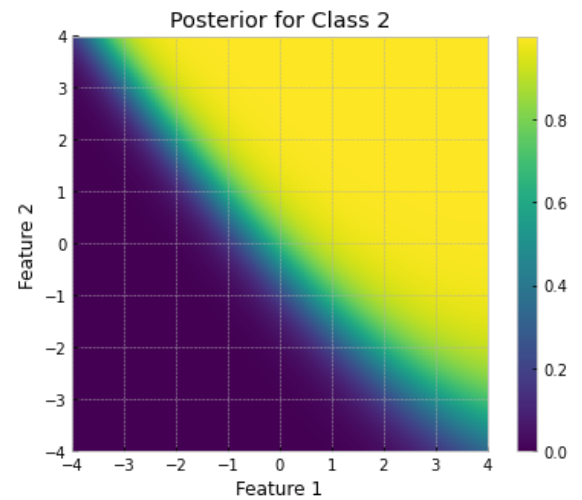
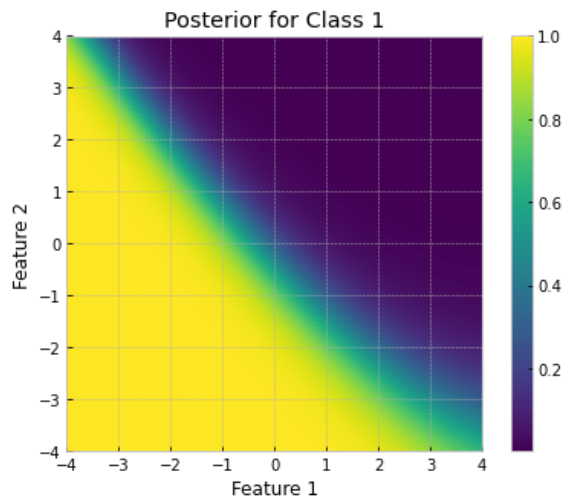
fig.add_subplot(1,2,2)
plt.scatter(data_C2[:,0], data_C2[:,1], c='b',alpha=0.3)
plt.imshow(y2, extent=[-4,4,-4,4], cmap='Blues')
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('PDF Likelihood for Class 2');
```



```
In [18]: # Let's take a look at the posterior distributions: they represent our classification d
pos1 = (y1*pC1)/(y1*pC1 + y2*pC2) # P(C1|x)
pos2 = (y2*pC2)/(y1*pC1 + y2*pC2) # P(C2|x)

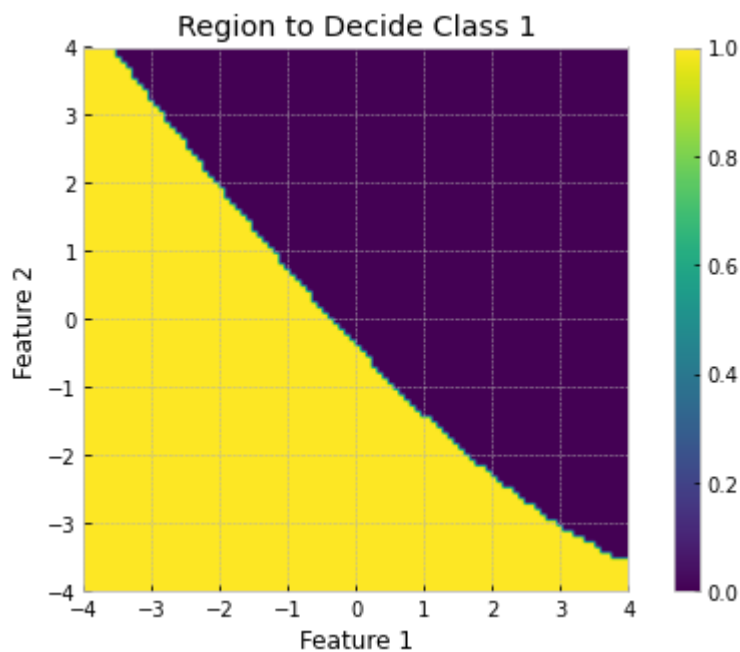
fig=plt.figure(figsize=(15,5))
fig.add_subplot(1,2,1)
plt.imshow(pos1, extent=[-4,4,-4,4])
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('Posterior for Class 1')

fig.add_subplot(1,2,2)
plt.imshow(pos2, extent=[-4,4,-4,4])
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('Posterior for Class 2');
```



In [19]:

```
# Look at the decision boundary:
plt.figure(figsize=(8,5))
plt.imshow(pos1 > pos2, extent=[-4,4,-4,4])
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('Region to Decide Class 1');
```



- Let's use this classifier to predict the class label for point [1, 1]:

In [21]:

```
x = [1,1]

# Data Likelihoods
y1_newPoint = multivariate_normal.pdf(x, mean=mu1, cov=cov1) # p(x|C1)
y2_newPoint = multivariate_normal.pdf(x, mean=mu2, cov=cov2) # p(x|C2)

print('Data likelihoods:')
print('P(x|C1) = ', y1_newPoint)
print('P(x|C2) = ', y2_newPoint, '\n')

# Posterior Probabilities
```



```

y1_pos = y1_newPoint*pC1/(y1_newPoint*pC1 + y2_newPoint*pC2)
y2_pos = y2_newPoint*pC2/(y1_newPoint*pC1 + y2_newPoint*pC2)

print('Posterior probabilities:')
print('P(C1|x) = ', y1_pos)
print('P(C2|x) = ', y2_pos, '\n')

if y1_pos > y2_pos:
    print('x = ',x,' belongs to class 1')
else:
    print('x = ',x,' belongs to class 2')

```

Data likelihoods:

$P(x|C1) = 0.003983945876533317$

$P(x|C2) = 0.16833065178108447$

Posterior probabilities:

$P(C1|x) = 0.011695292633899514$

$P(C2|x) = 0.9883047073661004$

$x = [1, 1]$  belongs to class 2

- What about  $x = [2, 4]$ ?

In [23]:

```

x = [2,4]

# Data Likelihoods
y1_newPoint = multivariate_normal.pdf(x, mean=mu1, cov=cov1) # p(x/C1)
y2_newPoint = multivariate_normal.pdf(x, mean=mu2, cov=cov2) # p(x/C2)

print('Data likelihoods:')
print('P(x|C1) = ', y1_newPoint)
print('P(x|C2) = ', y2_newPoint, '\n')

# Posterior Probabilities
y1_pos = y1_newPoint*pC1/(y1_newPoint*pC1 + y2_newPoint*pC2)
y2_pos = y2_newPoint*pC2/(y1_newPoint*pC1 + y2_newPoint*pC2)

print('Posterior probabilities:')
print('P(C1|x) = ', y1_pos)
print('P(C2|x) = ', y2_pos, '\n')

if y1_pos > y2_pos:
    print('x = ',x,' belongs to class 1')
else:
    print('x = ',x,' belongs to class 2')

```

Data likelihoods:

$P(x|C1) = 8.676399018565161e-08$

$P(x|C2) = 0.000872588499860459$

Posterior probabilities:

$P(C1|x) = 4.971396988221305e-05$

$P(C2|x) = 0.9999502860301178$

$x = [2, 4]$  belongs to class 2

Let's expand the decision surface view:

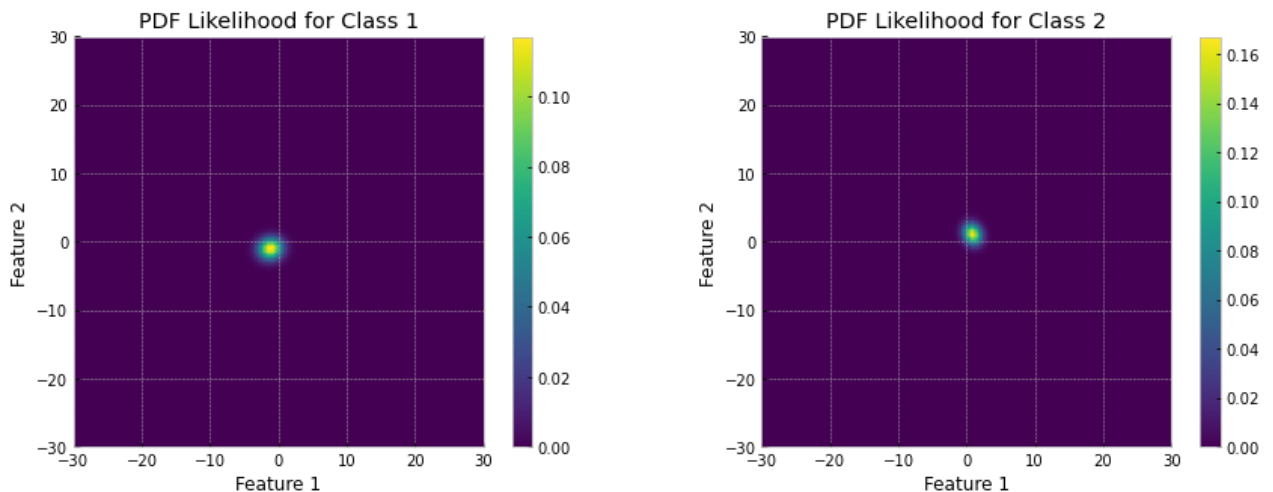
In [41]: `nd=30`

```
# Compute a grid of values for x and y
x = np.linspace(-nd, nd, 100)
y = np.linspace(-nd, nd, 100)
xm, ym = np.meshgrid(x, y)
X = np.flip(np.dstack([xm,ym]),axis=0)

# Let's plot the probabaility density function (pdf) for each class
y1 = multivariate_normal.pdf(X, mean=mu1, cov=cov1) #  $P(x|C1)$ 
y2 = multivariate_normal.pdf(X, mean=mu2, cov=cov2) #  $P(x|C2)$ 

fig =plt.figure(figsize=(15,5))
fig.add_subplot(1,2,1)
plt.imshow(y1, extent=[-nd,nd,-nd,nd])
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('PDF Likelihood for Class 1')

fig.add_subplot(1,2,2)
plt.imshow(y2, extent=[-nd,nd,-nd,nd])
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('PDF Likelihood for Class 2');
```

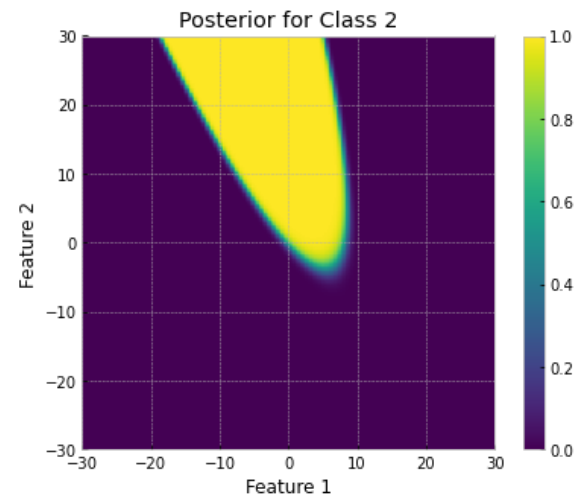
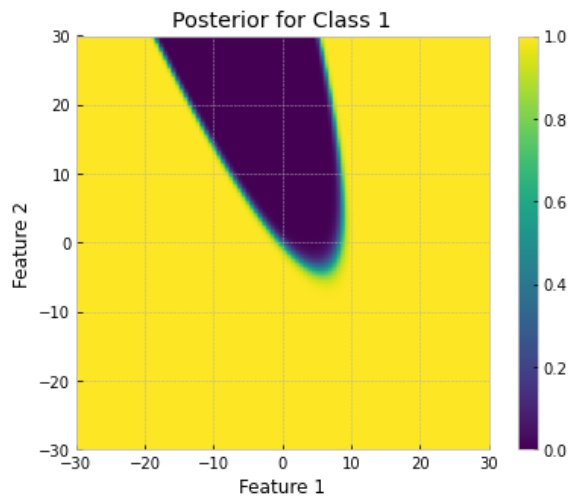


In [42]:

```
# Let's take a look at the posterior distributions: they represent our classification d
pos1 = (y1*pC1)/(y1*pC1 + y2*pC2) #  $P(C1|x)$ 
pos2 = (y2*pC2)/(y1*pC1 + y2*pC2) #  $P(C2|x)$ 

fig =plt.figure(figsize=(15,5))
fig.add_subplot(1,2,1)
plt.imshow(pos1, extent=[-nd,nd,-nd,nd])
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('Posterior for Class 1')

fig.add_subplot(1,2,2)
plt.imshow(pos2, extent=[-nd,nd,-nd,nd])
plt.colorbar()
plt.xlabel('Feature 1'); plt.ylabel('Feature 2')
plt.title('Posterior for Class 2');
```



- What about  $x = [10, 10]$ ?

In [43]:

```
x = [10,10]

# Data Likelihoods
y1_newPoint = multivariate_normal.pdf(x, mean=mu1, cov=cov1) #  $p(x|C1)$ 
y2_newPoint = multivariate_normal.pdf(x, mean=mu2, cov=cov2) #  $p(x|C2)$ 

print('Data likelihoods:')
print('P(x|C1) = ', y1_newPoint)
print('P(x|C2) = ', y2_newPoint, '\n')

# Posterior Probabilities
y1_pos = y1_newPoint*pC1/(y1_newPoint*pC1 + y2_newPoint*pC2)
y2_pos = y2_newPoint*pC2/(y1_newPoint*pC1 + y2_newPoint*pC2)

print('Posterior probabilities:')
print('P(C1|x) = ', y1_pos)
print('P(C2|x) = ', y2_pos, '\n')

if y1_pos > y2_pos:
    print('x = ',x,' belongs to class 1')
else:
    print('x = ',x,' belongs to class 2')
```

```
Data likelihoods:
P(x|C1) = 1.640177424643123e-41
P(x|C2) = 7.103322349039243e-45
```

```
Posterior probabilities:
P(C1|x) = 0.9991345844764138
P(C2|x) = 0.0008654155235862463
```

```
x = [10, 10] belongs to class 1
```

## Mixture Models

In this example, we look at a relatively simple model where we model each class with a single Gaussian probability density function (pdf).

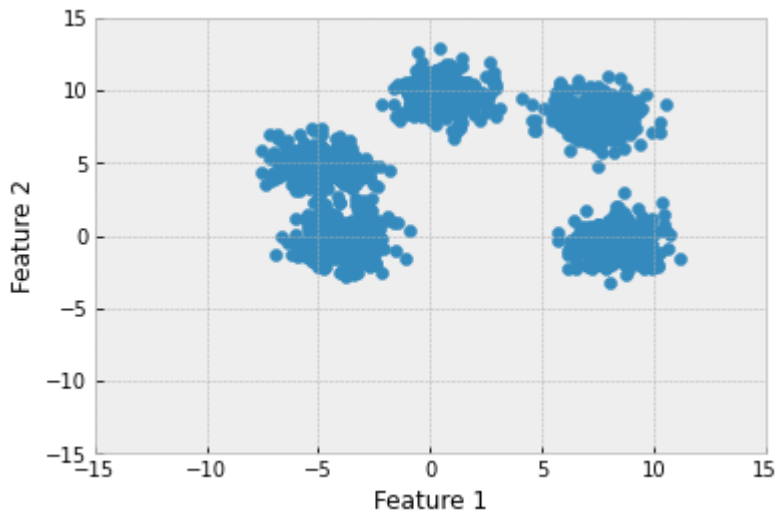
- What if the data for a single class looked like the plot below?

In [26]:

```
from sklearn.datasets import make_blobs

data, _ = make_blobs(n_samples = 1500, centers = 5)

plt.scatter(data[:,0],data[:,1]); plt.axis([-15,15,-15,15])
plt.xlabel('Feature 1'); plt.ylabel('Feature 2');
```



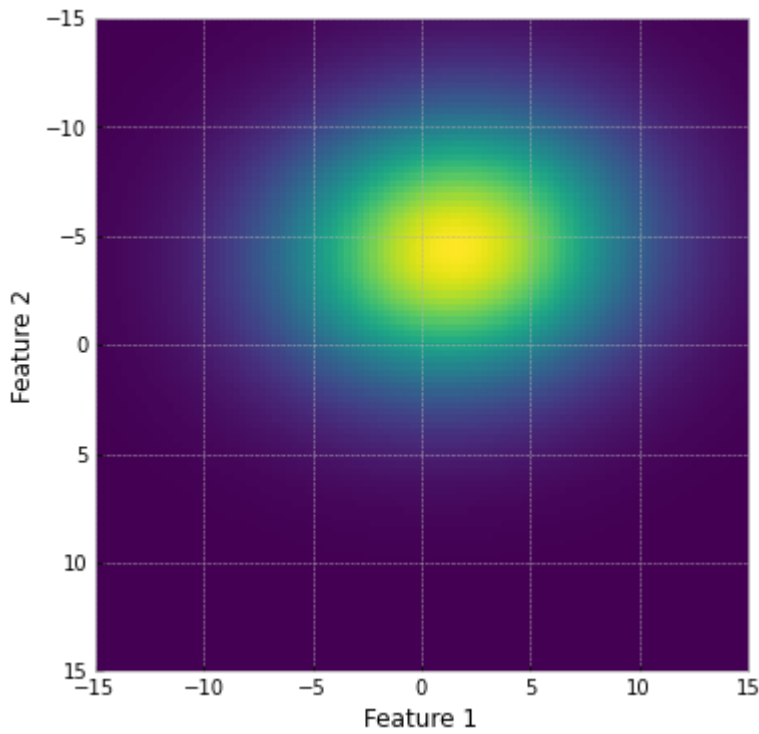
If we assume a single Gaussian distribution, we would obtain a very poor estimate of the true underlying data likelihood:

In [27]:

```
# Compute a grid of values for x and y
x = np.linspace(-15, 15, 100)
y = np.linspace(-15,15, 100)
xm, ym = np.meshgrid(x, y)
X = np.flip(np.dstack([xm,ym]),axis=0)

y1 = multivariate_normal.pdf(X, mean=np.mean(data, axis=0), cov=np.cov(data.T))

plt.figure(figsize=(6,6))
plt.imshow(y1, extent=[-15,15,15,-15])
plt.xlabel('Feature 1'); plt.ylabel('Feature 2');
```



## Mixture Models

We can better represent this data with a **mixture model**:

$$p(x|\Theta) = \sum_{k=1}^K \pi_k P(x|\Theta_k)$$

where  $\Theta = \{\Theta_k\}_{k=1}^K$  are set of parameters that define the distributional form in the probabilistic model  $P(\bullet|\Theta_k)$  and

$$0 \leq \pi_k \leq 1 \quad (19)$$

$$\sum_k \pi_k = 1 \quad (20)$$

If the probabilistic model  $P(\bullet|\Theta_k)$  is assumed to be a Gaussian distribution, then the above mixture model is a **Gaussian Mixture Model (GMM)**

$$P(x|\Theta) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$$

where  $\Theta_k = \{\mu_k, \Sigma_k, \pi_k\}, \forall k$  are the mean, covariance and weight of each Gaussian distribution, and, again,  $0 \leq \pi_k \leq 1$  with  $\sum_k \pi_k = 1$ .