


# Lecture 33 - Slack Variables; Soft-Margin SVM

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('bmh')
```


## Example of Hard-Margin SVM

Linearly-separable classes in the feature space spanned by the transformation  $\phi(x)$  may form a non-linear decision boundary in the input space.

```
In [2]: from IPython.display import Image
Image("figures/Figure7.2.png", width=400)
```



```
In [3]: from sklearn.datasets import make_blobs
X, y = make_blobs(n_samples=50, centers=2,
                  random_state=0, cluster_std=0.60)
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn');
```



```
In [4]: from sklearn.svm import SVC # "Support vector classifier"

model = SVC(kernel='linear')
model.fit(X, y)
```

```
Out[4]: SVC(kernel='linear')
```

```
In [5]: # Python Data Science Handbook


def plot_svc_decision_function(model, ax=None, plot_support=True):
    """Plot the decision function for a 2D SVC"""
    if ax is None:
        ax = plt.gca()
    xlim = ax.get_xlim()
    ylim = ax.get_ylim()

    # create grid to evaluate model
    x = np.linspace(xlim[0], xlim[1], 30)
    y = np.linspace(ylim[0], ylim[1], 30)
    Y, X = np.meshgrid(y, x)
    xy = np.vstack([X.ravel(), Y.ravel()]).T
    P = model.decision_function(xy).reshape(X.shape)

    # plot decision boundary and margins
    ax.contour(X, Y, P, colors='k',
               levels=[-1, 0, 1], alpha=0.5,
               linestyles=['--', '-', '--'])

    # plot support vectors
    if plot_support:
        ax.scatter(model.support_vectors_[:, 0],
                   model.support_vectors_[:, 1],
                   s=300, linewidth=1, edgecolors='black', facecolors='none');
    ax.set_xlim(xlim)
    ax.set_ylim(ylim)
```

```
In [6]: plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
plot_svc_decision_function(model);
```



```
In [7]: model.support_vectors_
```

```
Out[7]: array([[0.44359863, 3.11530945],
               [2.33812285, 3.43116792],
               [2.06156753, 1.96918596]])
```

```
In [8]: def plot_svm(N=10, ax=None):
        X, y = make_blobs(n_samples=200, centers=2,
                           random_state=0, cluster_std=0.60)

        X = X[:N]
        y = y[:N]
        model = SVC(kernel='linear', C=1E10)
        model.fit(X, y)


        ax = ax or plt.gca()
        ax.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
        ax.set_xlim(-1, 4)
        ax.set_ylim(-1, 6)
        plot_svc_decision_function(model, ax)

        from ipywidgets import interact, fixed
        interact(plot_svm, N=[10, 30, 60, 100, 200], ax=fixed(None));
```

```
In [9]: from sklearn.datasets import make_circles
        X, y = make_circles(100, factor=.1, noise=.1)

        clf = SVC(kernel='linear').fit(X, y)


        plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
        plot_svc_decision_function(clf, plot_support=False);
```



```
In [10]: clf = SVC(kernel='rbf', C=1E6)
         clf.fit(X, y)
```

```
Out[10]: SVC(C=1000000.0)
```

```
In [11]: plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
         plot_svc_decision_function(clf)
         plt.scatter(clf.support_vectors_[:, 0], clf.support_vectors_[:, 1],
                    s=300, lw=1, facecolors='none');
```



## Soft-Margin Support Vector Machine (SVM): Overlapping Classes

To handle this case, the SVM implementation has a bit of a fudge-factor which "softens" the margin: that is, it allows some of the points to creep into the margin if that allows a better fit. The hardness of the margin is controlled by a tuning parameter, most often known as **slack variable**  $\xi_n \geq 0$ ,  $n = 1, \dots, N$ , with one slack variable for each training data point. For very large  $\xi$ , the margin is hard, and points cannot lie in it. For smaller  $\xi$ , the margin is softer, and can grow to encompass some points.

A **slack variable** is defined as  $\xi_n = 0$  for data points that are on or inside the correct margin boundary and  $\xi_n = |t_n - y(x_n)|$  for other points. Thus a data point that is on the decision boundary  $y(x_n) = 0$  will have  $\xi_n = 1$ , and points with  $\xi_n > 1$  will be misclassified. The exact classification constraints are then replaced with


$$t_n y(x_n) \geq 1 - \xi_n, n = 1, \dots, N$$

in which the slack variables are constrained to satisfy  $\xi_n \geq 0$ .

- Data points for which  $\xi_n = 0$  are correctly classified and are either on the margin or on the correct side of the margin.
- Points for which  $0 < \xi_n \leq 1$  lie inside the margin, but on the correct side of the decision boundary.
- And those data points for which  $\xi_n > 1$  lie on the wrong side of the decision boundary and are misclassified.

```
In [12]: Image('figures/Figure7.3.png', width=400)
```

Out[12]:



Our goal is now to maximize the margin while softly penalizing points that lie on the wrong side of the margin boundary. We therefore minimize:

$$\begin{aligned} \arg_{w,b} \min C \sum_{n=1}^N \xi_n + \frac{1}{2} \|w\|^2 \\ \text{subject to } t_n y(x_n) \geq 1 - \xi_n, n = 1, \dots, N \\ \text{and } \xi_n \geq 0, n = 1, \dots, N \end{aligned}$$

where the parameter  $C > 0$  controls the trade-off between the slack variable penalty and the margin.

- Because any point that is misclassified has  $\xi_n > 1$ , it follows that  $\sum_n \xi_n$  is an upper bound on the number of misclassified points.

- The parameter  $C$  is therefore analogous to (the inverse of) a regularization coefficient because it controls the trade-off between minimizing training errors and controlling model complexity.
- In the limit  $C \rightarrow \infty$ , we will recover the earlier support vector machine for separable data.

The Lagrangian is given by:

$$L(w, b, a) = \frac{1}{2} \|w\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N a_n (t_n y(x_n) - 1 + \xi_n) - \sum_{n=1}^N \mu_n \xi_n$$

where  $\{a_n \geq 0\}_{n=1}^N$  and  $\{\mu_n \geq 0\}_{n=1}^N$  are Lagrange multipliers. The corresponding set of Karush–Kuhn–Tucker (KKT) conditions are given by

$$\begin{aligned} a_n &\geq 0 \\ t_n y(x_n) - 1 + \xi_n &\geq 0 \\ a_n (t_n y(x_n) - 1 + \xi_n) &= 0 \\ \mu_n &\geq 0 \\ \xi_n &\geq 0 \\ \mu_n \xi_n &= 0 \end{aligned}$$

where  $n = 1, \dots, N$ .

We now optimize for  $w$ ,  $b$  and  $\{\xi_n\}$ :

$$\begin{aligned} \frac{\partial L}{\partial w} = 0 &\Rightarrow w = \sum_{n=1}^N a_n t_n \phi(x_n) \\ \frac{\partial L}{\partial b} = 0 &\Rightarrow \sum_{n=1}^N a_n t_n = 0 \\ \frac{\partial L}{\partial \xi_n} = 0 &\Rightarrow a_n = C - \mu_n \Rightarrow a_n \leq C \end{aligned}$$

The dual Lagrangian is then given by:

$$\tilde{L}(a) = \sum_{n=1}^N a_n - \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m k(x_n, x_m)$$

which is identical to the separable case, except that the constraints are somewhat different. We therefore have to minimize  $\tilde{L}(a)$  with respect to the dual variables  $\{a_n\}$  subject to

$$0 \leq a_n \leq C$$

$$\sum_{n=1}^N a_n t_n = 0$$

As before, a subset of the data points may have  $a_n = 0$ , in which case they do not contribute to the predictive model. The remaining data points constitute the support vectors. These have  $a_n > 0$  and hence  $t_n y(x_n) = 1 - \xi_n$ .

- If  $a_n < C$ , then  $\mu_n > 0$ , which requires  $\xi_n = 0$  and hence such points lie on the margin.

- Points with  $a_n = C$  can lie inside the margin and can either be correctly classified if  $\xi_n \leq 1$  or misclassified if  $\xi_n > 1$ .

To determine the parameter  $b$ , we note that those support vectors for which  $0 < a_n < C$  have  $\xi_n = 0$  so that  $t_n y(x_n) = 1$  and hence will satisfy

$$t_n \left( \sum_{m \in S} a_m t_m k(x_n, x_m) + b \right) = 1$$

Again, a numerically stable solution is obtained by averaging to give

$$b = \frac{1}{N_M} \sum_{n \in M} \left( t_n - \sum_{m \in S} a_m t_m k(x_n, x_m) \right)$$

where  $M$  denotes the set of indices of data points having  $0 < a_n < C$ .

Although predictions for new inputs are made using only the support vectors, the training phase (i.e., the determination of the parameters  $a$  and  $b$ ) makes use of the whole data set, and so it is important to have **efficient algorithms for solving the quadratic programming problem**.

We first note that the objective function  $\tilde{L}(a)$  is quadratic and so any local optimum will also be a **global optimum** provided the constraints define a convex region (which they do as a consequence of being linear).